

Meridian Product Roadmap 2026

Meridian Global Corp — Confidential Internal Document

Meridian Product Roadmap 2026

Document ID: MGC-ROADMAP-2026-01

Effective Date: January 15, 2026

Classification: Internal — Executive & Engineering Leadership

Last Updated: January 10, 2026

Executive Summary

Meridian Global Corp's 2026 Product Roadmap represents a pivotal year of platform transformation, customer-led innovation, and operational hardening. With 40,000 employees serving over 8,500 enterprise clients across 120 countries, we are committing our engineering resources to four strategic pillars: AI/ML capabilities (40% of total engineering capacity), platform reliability (30%), enterprise features (20%), and developer experience (10%). The roadmap is structured in four quarterly phases, each delivering measurable outcomes that align with our mission to make Meridian the most intelligent, secure, and extensible technology platform for global enterprises.

Q1 2026 — Foundation Modernization & Developer Reach

GraphQL API (General Availability)

Effective February 3, 2026, Meridian's GraphQL API moves from beta to GA. This release eliminates over-prefetching and enables declarative data queries. Customers migrating from REST will receive a 40% reduction in payload size and a 55% improvement in response latency under load. The GA release includes a dedicated SDK for Python 3.12, Node.js 22, and Java 21, plus interactive documentation via Apollo Studio.

Single Sign-On (SAML 2.0) Integration

On March 17, 2026, we launch native SAML 2.0 federation supporting Identity Providers (Okta, Azure AD, OneLogin, PingIdentity). This covers IdP-initiated and SP-initiated flows, just-in-time provisioning, and attribute mapping for role-based access groups. Existing OIDC users can coexist; administrators may enable both protocols per tenant.

****Mobile App Beta****

The Meridian Mobile Beta (iOS 17+ and Android 14+) launches March 31, 2026. Initial features include push notification management, approval workflows, real-time dashboard views, and biometric authentication. Beta testers will be recruited from 200 existing enterprise accounts; full public release is scheduled for Q3 2026.

Q2 2026 — Predictive Intelligence & Enterprise Depth

****Advanced Analytics — Predictive Models (GA)****

On April 20, 2026, we release the Insights Predictive Engine. Built on a custom XGBoost time-series architecture, it delivers anomaly detection, forecast generation (inventory, churn, capacity), and natural-language explanations of model outputs. Customers can train models on up to 7 years of historical data with no additional infrastructure cost. SLA: 99.5% uptime for model serving endpoints.

****Custom Embedding Support****

May 11, 2026 introduces a developer portal for custom vector embeddings. Customers may inject domain-specific embeddings (e.g., financial risk scores, medical ontologies) into Meridian's search and recommendation pipelines. API endpoints support batching of up to 10,000 vectors per request with sub-100ms latency for real-time queries.

****Audit Logging (Enhanced)****

Starting June 1, 2026, all Meridian SaaS tenants receive immutable, JSON-formatted audit logs with a retention period configurable from 90 to 1,095 days. Logs capture every API call, configuration change, authentication event, and data access. Integration with Splunk, Datadog, and Sumo Logic is prebuilt. Compliance reports (SOC 2 Type II, ISO 27001) will reference this feature from Q3 2026 onward.

Q3 2026 — Meridian 3.0 Launch: AI-Native Architecture

****Meridian 3.0 Platform Release****

On August 3, 2026, Meridian 3.0 becomes available. This is a ground-up rearchitecture replacing the legacy microservices mesh with an AI-native orchestration layer that dynamically allocates compute, memory, and GPU resources using a neural scheduler. Key metrics:

- ****10x performance improvement**** in data pipeline throughput (measured from ingestion to analytics output) under 50 concurrent tenants.
- ****Latency reduction****: 95th-percentile API response time drops from 320ms to 45ms.
- ****On-premise option****: A fully air-gapped deployment for regulated industries (defense, healthcare, finance) using Kubernetes 1.28 and RHEL 9.4. Licensed per node with a 24-node minimum.

Migration assistance includes automated pre-checks, zero-downtime cutover tools for SaaS customers, and a dedicated migration engineering team. On-premise systems will reach GA on September 14, 2026.

Q4 2026 — Ecosystem Scale & Compliance Hardening

****Meridian Marketplace (100+ Connectors)****

On October 26, 2026, the Meridian Marketplace opens with 117 prebuilt connectors spanning CRM (Salesforce, HubSpot), ERP (SAP S/4HANA, Oracle NetSuite), data warehouses (Snowflake, BigQuery, Redshift), productivity tools (Slack, Teams, Jira), and legacy systems (AS/400 via MQ). Connectors support bi-directional sync, field-level mapping, and scheduling. Third-party developers can publish connectors using a certified SDK; revenue share is 70/30 in favor of the developer.

****Multi-Tenant SaaS Hardening****

November 9, 2026 sees the completion of a 7-month tenant isolation project. Each tenant now runs in a dedicated eBPF-based sidecar with memory encryption (AMD SEV-SNP) and per-tenant rate limiting. Tenant boot time reduces from 4.2 seconds to 0.8 seconds. A new tenant can be provisioned in under 90 seconds, meeting SLA for automatic scaling during acquisition events.

****HIPAA Compliance (Covered Entity)****

Effective December 1, 2026, Meridian's SaaS and on-premise 3.0 platforms are HIPAA-compliant for covered entities. A Business Associate Agreement (BAA) is executed upon request. Features include auditable PHI access controls, automatic encryption at rest (AES-256-GCM) and in transit (TLS 1.3), ePHI logging with 6-year retention, and breach notification automation. Third-party certification via HITRUST is expected by January 2027.

Investment Priorities (Engineering Allocation)

Priority Area	% of Engineering Capacity	Rationale
---------------	---------------------------	-----------

---	---	---
-----	-----	-----

AI/ML Capabilities	40%	Predictive analytics, AI-native scheduler, embedding
--------------------	-----	--